

# Lab 3: Gesture Recognition using Convolutional Neural Networks

In this lab you will train a convolutional neural network to make classifications on different hand gestures. By the end of the lab, you should be able to:

1. Load and split data for training, validation and testing
2. Train a Convolutional Neural Network
3. Apply transfer learning to improve your model

Note that for this lab we will not be providing you with any starter code. You should be able to take the code used in previous labs, tutorials and lectures and modify it accordingly to complete the tasks outlined below.

## What to submit

Submit a PDF file containing all your code, outputs, and write-up from parts 1-5. You can produce a PDF of your Google Colab file by going to **File > Print** and then save as PDF. The Colab instructions has more information. Make sure to review the PDF submission to ensure that your answers are easy to read. Make sure that your text is not cut off at the margins.

**Do not submit any other files produced by your code.**

Include a link to your colab file in your submission.

Please use Google Colab to complete this assignment. If you want to use Jupyter Notebook, please complete the assignment and upload your Jupyter Notebook file to Google Colab for submission.

## Colab Link

Include a link to your colab file here

Colab Link: <https://drive.google.com/file/d/1JClA5Ge0TNi2XObKjcJuHFodIEcWdpJe/view?usp=sharing>

## Dataset

American Sign Language (ASL) is a complete, complex language that employs signs made by moving the hands combined with facial expressions and postures of the body. It is the primary language of many North Americans who are deaf and is one of several communication options used by people who are deaf or hard-of-hearing. The hand gestures representing English alphabet are shown below. This lab focuses on classifying a subset of these hand gesture images using convolutional neural networks. Specifically, given an image of a hand showing one of the letters A-I, we want to detect which letter is being represented.



## Part B. Building a CNN [50 pt]

For this lab, we are not going to give you any starter code. You will be writing a convolutional neural network from scratch. You are welcome to use any code from previous labs, lectures and tutorials. You should also write your own code.

You may use the PyTorch documentation freely. You might also find online tutorials helpful. However, all code that you submit must be your own.

Make sure that your code is vectorized, and does not contain obvious inefficiencies (for example, unnecessary for loops, or unnecessary calls to `unsqueeze()`). Ensure enough comments are included in the code so that your TA can understand what you are doing. It is your responsibility to show that you understand what you write.

**This is much more challenging and time-consuming than the previous labs.** Make sure that you give yourself plenty of time by starting early.

### 1. Data Loading and Splitting [5 pt]

Download the anonymized data provided on Quercus. To allow you to get a heads start on this project we will provide you with sample data from previous years. Split the data into training, validation, and test sets.

Note: Data splitting is not as trivial in this lab. We want our test set to closely resemble the setting in which our model will be used. In particular, our test set should contain hands that are never seen in training!

Explain how you split the data, either by describing what you did, or by showing the code that you used. Justify your choice of splitting strategy. How many training, validation, and test images do you have?

For loading the data, you can use `plt.imread` as in Lab 1, or any other method that you choose. You may find `torchvision.datasets.ImageFolder` helpful. (see <https://pytorch.org/docs/stable/torchvision/datasets.html?highlight=image%20folder#torchvision.datasets.ImageFolder> )

```
In [1]: import numpy as np
import time
import torch
import torch.nn as nn
import torch.nn.functional as F
import torch.optim as optim
import torchvision
from torchvision import transforms, datasets

torch.manual_seed(1)
```

```
Out[1]: <torch._C.Generator at 0x1e7d9696b70>
```

```
In [2]: # Load entire dataset (not split)
file_dir = "C:/Users/Admin/Downloads/Lab3_Dataset/Lab3_Gestures_Summer"
transform = transforms.Compose([
    transforms.Resize((224, 224)),
    transforms.ToTensor(),          # Convert to tensor
])

dataset = datasets.ImageFolder(root=file_dir, transform=transform)
```

```
In [3]: classes = ['A', 'B', 'C', 'D', 'E', 'F', 'G', 'H', 'I']
```

```
In [4]: def get_data_loader(batch_size, overfitting=False):
    """ Splits the dataset into training, validation and testing datasets.
    Returns data loaders for the three preprocessed datasets.

    Args:
        batch_size: A int representing the number of samples per batch

    Returns:
        train_loader: iterable training dataset organized according to batch size
        val_loader: iterable validation dataset organized according to batch size
        test_loader: iterable testing dataset organized according to batch size
    """
    # Define size
    train_size = int(0.8 * len(dataset))
    val_size = int(0.1 * len(dataset))
    test_size = len(dataset) - train_size - val_size
    overfit_size = 0
    if (overfitting):
        val_size = int(0.09 * len(dataset))
        test_size = int(0.09 * len(dataset))
        overfit_size = len(dataset) - train_size - val_size - test_size

    # Random split
    train_dataset, val_dataset, test_dataset, overfit_dataset = torch.utils.data.random_split(
        dataset, [train_size, val_size, test_size, overfit_size]
    )

    # Create data loaders for each set
    train_loader = torch.utils.data.DataLoader(train_dataset, batch_size=batch_size)
    val_loader = torch.utils.data.DataLoader(val_dataset, batch_size=batch_size)
```

```
test_loader = torch.utils.data.DataLoader(test_dataset, batch_size=batch_size)
overfit_loader = torch.utils.data.DataLoader(overfit_dataset, batch_size=batch_size)

return train_loader, val_loader, test_loader, overfit_loader
```

```
In [5]: ## TEST visualize dataset
import matplotlib.pyplot as plt

train_loader, val_loader, test_loader, overfit_loader = get_data_loader(batch_size = 1)

k = 0
for images, labels in train_loader:
    plt.subplot(3, 5, k+1)
    plt.axis('off')
    plt.imshow(np.transpose(images[0], (1, 2, 0)))

    k += 1
    if k > 14:
        break
```



```
In [6]: print("Training data size: ", len(train_loader) * train_loader.batch_size)
print("Validation data size: ", len(val_loader) * val_loader.batch_size)
print("Testing data size: ", len(test_loader) * test_loader.batch_size)
print("Overfitting data size: ", len(overfit_loader) * test_loader.batch_size)
```

```
Training data size: 1775
Validation data size: 221
Testing data size: 223
Overfitting data size: 0
```

## 2. Model Building and Sanity Checking [15 pt]

### Part (a) Convolutional Network - 5 pt

Build a convolutional neural network model that takes the (224x224 RGB) image as input, and predicts the gesture letter. Your model should be a subclass of `nn.Module`. Explain your choice of neural network architecture: how many layers did you choose? What types of layers did you use? Were they fully-connected or convolutional? What about other decisions like pooling layers, activation functions, number of channels / hidden units?

Architecture:

**3x224x224 (conv1) 5x220x220 (maxpool) 5x110x110 (conv2) 10x106x106 (pool) 10x53x53 (flattened, relu, fc1) 244 (fc2) 9 (Output)**

- CNN, then feed to FC (good for image classification)
- ReLU activation function (recommended in lecture)
- Pooling: Max pooling (better than average pooling, from lecture)
- Number of channels & kernel size (based on lab 2 - Seems ok for Google Colab computation power)

```
In [7]: class Gesture(nn.Module):
    def __init__(self, name):
        super(Gesture, self).__init__()
        self.name = name
        self.conv1 = nn.Conv2d(3, 5, 5)
        self.pool = nn.MaxPool2d(2, 2)
        self.conv2 = nn.Conv2d(5, 10, 5)
        self.fc1 = nn.Linear(10 * 53 * 53, 244)
        self.fc2 = nn.Linear(244, 9) # ABCDEFGHI

    def forward(self, x):
        x = self.pool(F.relu(self.conv1(x)))
        x = self.pool(F.relu(self.conv2(x)))
        x = x.view(-1, 10 * 53 * 53)
        x = F.relu(self.fc1(x))
        x = self.fc2(x)
        return x
```

## Part (b) Training Code - 5 pt

Write code that trains your neural network given some training data. Your training code should make it easy to tweak the usual hyperparameters, like batch size, learning rate, and the model object itself. Make sure that you are checkpointing your models from time to time (the frequency is up to you). Explain your choice of loss function and optimizer.

```
In [39]: #####
# Training Function
def train(model, overfitting=False, batch_size=30, learning_rate=0.001, num_epochs=30, replace=[
    #####
    # Fixed PyTorch random seed for reproducible result
    torch.manual_seed(100)
    #####
    # Obtain the PyTorch data loader objects to load batches of the datasets
    train_loader, val_loader, test_loader, overfit_loader = get_data_loader(batch_size, overfitt
    training = train_loader
    validating = val_loader
    if (overfitting):
        print("WARNING: only overfit testing should be here")
        training = overfit_loader
        validating = overfit_loader
    if (len(replace) > 0):
        print("WARNING: only transfer learning should be here")
        training = replace[0]
        validating = replace[1]
    #####
    # Define the Loss function and optimizer
    # The loss function will be Cross Entropy (CE) - since multiple classes involved.
    # We will use the nn.CrossEntropyLoss() - input expects logit (un-normalized output)
    # Optimizer will be Adam (recommended in class)
    criterion = nn.CrossEntropyLoss()
    optimizer = optim.Adam(model.parameters(), lr=learning_rate)
    #####
    # Set up some numpy arrays to store the training/test loss/error
    train_err = np.zeros(num_epochs)
    train_loss = np.zeros(num_epochs)
```

```

val_err = np.zeros(num_epochs)
val_loss = np.zeros(num_epochs)
#####
# Train the network
# Loop over the data iterator and sample a new batch of training data
# Get the output from the network, and optimize our loss function.
for epoch in range(num_epochs): # Loop over the dataset multiple times
    total_train_loss = 0.0
    total_train_err = 0.0
    total_epoch = 0
    for i, data in enumerate(training, 0):
        # Get the inputs
        inputs, labels = data
        labels = F.one_hot(labels, len(classes))
        if torch.cuda.is_available():
            inputs = inputs.cuda()
            labels = labels.cuda()

        # Zero the parameter gradients
        optimizer.zero_grad()
        # Forward pass, backward pass, and optimize
        outputs = model(inputs)
        loss = criterion(outputs, labels.float())
        loss.backward()
        optimizer.step()
        # Calculate the statistics
        _, pred = outputs.max(dim=1)
        total_train_err += len(labels) - pred.eq(labels.argmax(dim=1)).sum().item()
        total_train_loss += loss.item()
        total_epoch += len(labels)
    train_err[epoch] = float(total_train_err) / total_epoch
    train_loss[epoch] = float(total_train_loss) / (i+1)
    val_err[epoch], val_loss[epoch] = evaluate(model, validating, criterion)
    print("Epoch {}: Train err: {}, Train loss: {} | "+
          "Validation err: {}, Validation loss: {}".format(
            epoch + 1,
            train_err[epoch],
            train_loss[epoch],
            val_err[epoch],
            val_loss[epoch]))
    # Save the current model (checkpoint) to a file
    model_path = get_model_name(model.name, batch_size, learning_rate, epoch)
    torch.save(model.state_dict(), model_path)
print('Finished Training')
# Write the train/test loss/err into CSV file for plotting later
epochs = np.arange(1, num_epochs + 1)
np.savetxt("{}_train_err.csv".format(model_path), train_err)
np.savetxt("{}_train_loss.csv".format(model_path), train_loss)
np.savetxt("{}_val_err.csv".format(model_path), val_err)
np.savetxt("{}_val_loss.csv".format(model_path), val_loss)

#####
# Helper Functions
def evaluate(model, loader, criterion):
    """ Evaluate the network on the validation set.

    Args:
        model: PyTorch neural network object
        loader: PyTorch data loader for the validation set
        criterion: The loss function

    Returns:
        err: A scalar for the avg classification error over the validation set
        loss: A scalar for the average loss function over the validation set
    """

```

```

total_loss = 0.0
total_err = 0.0
total_epoch = 0
for i, data in enumerate(loader, 0):
    inputs, labels = data
    labels = F.one_hot(labels, len(classes))
    if torch.cuda.is_available():
        inputs = inputs.cuda()
        labels = labels.cuda()

    outputs = model(inputs)

    ##### Statistics #####
    loss = criterion(outputs, labels.float()) # CE expects logits
    _, pred = outputs.max(dim=1)
    total_err += len(labels) - pred.eq(labels.argmax(dim=1)).sum().item()
    total_loss += loss.item()
    total_epoch += len(labels)
err = float(total_err) / total_epoch
loss = float(total_loss) / (i + 1)
return err, loss

def get_model_name(name, batch_size, learning_rate, epoch):
    """ Generate a name for the model consisting of all the hyperparameter values

    Args:
        config: Configuration object containing the hyperparameters
    Returns:
        path: A string with the hyperparameter name and value concatenated
    """
    path = "C:/Users/Admin/Downloads/lab3points/model_{0}_bs{1}_lr{2}_epoch{3}".format(name,
                                                                                       batch_size,
                                                                                       learning_rate,
                                                                                       epoch)

    return path

#####
# Training Curve
def plot_training_curve(path):
    """ Plots the training curve for a model run, given the csv files
    containing the train/validation error/loss.

    Args:
        path: The base path of the csv files produced during training
    """
    import matplotlib.pyplot as plt
    train_err = np.loadtxt("{}_train_err.csv".format(path))
    val_err = np.loadtxt("{}_val_err.csv".format(path))
    train_loss = np.loadtxt("{}_train_loss.csv".format(path))
    val_loss = np.loadtxt("{}_val_loss.csv".format(path))
    plt.title("Train vs Validation Error")
    n = len(train_err) # number of epochs
    plt.plot(range(1,n+1), train_err, label="Train")
    plt.plot(range(1,n+1), val_err, label="Validation")
    plt.xlabel("Epoch")
    plt.ylabel("Error")
    plt.legend(loc='best')
    plt.show()
    plt.title("Train vs Validation Loss")
    plt.plot(range(1,n+1), train_loss, label="Train")
    plt.plot(range(1,n+1), val_loss, label="Validation")
    plt.xlabel("Epoch")
    plt.ylabel("Loss")

```

```
plt.legend(loc='best')  
plt.show()
```

## Part (c) “Overfit” to a Small Dataset - 5 pt

One way to sanity check our neural network model and training code is to check whether the model is capable of “overfitting” or “memorizing” a small dataset. A properly constructed CNN with correct training code should be able to memorize the answers to a small number of images quickly.

Construct a small dataset (e.g. just the images that you have collected). Then show that your model and training code is capable of memorizing the labels of this small data set.

With a large batch size (e.g. the entire small dataset) and learning rate that is not too high, You should be able to obtain a 100% training accuracy on that small dataset relatively quickly (within 200 iterations).

```
In [9]: model = Gesture("Overfit")  
if torch.cuda.is_available():  
    model = model.cuda()  
train(model, overfitting=True, batch_size=46, learning_rate=0.001, num_epochs=30)
```



Epoch 1: Train err: 0.9347826086956522, Train loss: 2.208139657974243 | Validation err: 0.8043478260869565, Validation loss: 2.117039442062378

Epoch 2: Train err: 0.8043478260869565, Train loss: 2.117039442062378 | Validation err: 0.8043478260869565, Validation loss: 2.1230697631835938

Epoch 3: Train err: 0.8043478260869565, Train loss: 2.1230697631835938 | Validation err: 0.6739130434782609, Validation loss: 2.0661749839782715

Epoch 4: Train err: 0.6739130434782609, Train loss: 2.0661749839782715 | Validation err: 0.8043478260869565, Validation loss: 2.0482635498046875

Epoch 5: Train err: 0.8043478260869565, Train loss: 2.0482635498046875 | Validation err: 0.6521739130434783, Validation loss: 2.0252439975738525

Epoch 6: Train err: 0.6521739130434783, Train loss: 2.0252439975738525 | Validation err: 0.7608695652173914, Validation loss: 1.996256709098816

Epoch 7: Train err: 0.7608695652173914, Train loss: 1.996256709098816 | Validation err: 0.7608695652173914, Validation loss: 1.9547805786132812

Epoch 8: Train err: 0.7608695652173914, Train loss: 1.9547805786132812 | Validation err: 0.7608695652173914, Validation loss: 1.9224475622177124

Epoch 9: Train err: 0.7608695652173914, Train loss: 1.9224475622177124 | Validation err: 0.717391304347826, Validation loss: 1.8764595985412598

Epoch 10: Train err: 0.717391304347826, Train loss: 1.8764595985412598 | Validation err: 0.6956521739130435, Validation loss: 1.8280154466629028

Epoch 11: Train err: 0.6956521739130435, Train loss: 1.8280154466629028 | Validation err: 0.5, Validation loss: 1.7397921085357666

Epoch 12: Train err: 0.5, Train loss: 1.7397921085357666 | Validation err: 0.4782608695652174, Validation loss: 1.6828545331954956

Epoch 13: Train err: 0.4782608695652174, Train loss: 1.6828545331954956 | Validation err: 0.43478260869565216, Validation loss: 1.597893476486206

Epoch 14: Train err: 0.43478260869565216, Train loss: 1.597893476486206 | Validation err: 0.41304347826086957, Validation loss: 1.4992364645004272

Epoch 15: Train err: 0.41304347826086957, Train loss: 1.4992364645004272 | Validation err: 0.34782608695652173, Validation loss: 1.4072232246398926

Epoch 16: Train err: 0.34782608695652173, Train loss: 1.4072232246398926 | Validation err: 0.2608695652173913, Validation loss: 1.2776827812194824

Epoch 17: Train err: 0.2608695652173913, Train loss: 1.2776827812194824 | Validation err: 0.2608695652173913, Validation loss: 1.1478745937347412

Epoch 18: Train err: 0.2608695652173913, Train loss: 1.1478745937347412 | Validation err: 0.21739130434782608, Validation loss: 1.0486747026443481

Epoch 19: Train err: 0.21739130434782608, Train loss: 1.0486747026443481 | Validation err: 0.3695652173913043, Validation loss: 1.1473547220230103

Epoch 20: Train err: 0.3695652173913043, Train loss: 1.1473547220230103 | Validation err: 0.13043478260869565, Validation loss: 0.8395739197731018

Epoch 21: Train err: 0.13043478260869565, Train loss: 0.8395739197731018 | Validation err: 0.3695652173913043, Validation loss: 0.9892588257789612

Epoch 22: Train err: 0.3695652173913043, Train loss: 0.9892588257789612 | Validation err: 0.10869565217391304, Validation loss: 0.6576091051101685

Epoch 23: Train err: 0.10869565217391304, Train loss: 0.6576091051101685 | Validation err: 0.13043478260869565, Validation loss: 0.660605788230896

Epoch 24: Train err: 0.13043478260869565, Train loss: 0.660605788230896 | Validation err: 0.13043478260869565, Validation loss: 0.5606409311294556

Epoch 25: Train err: 0.13043478260869565, Train loss: 0.5606409311294556 | Validation err: 0.06521739130434782, Validation loss: 0.473480224609375

Epoch 26: Train err: 0.06521739130434782, Train loss: 0.473480224609375 | Validation err: 0.043478260869565216, Validation loss: 0.4257756173610687

Epoch 27: Train err: 0.043478260869565216, Train loss: 0.4257756173610687 | Validation err: 0.0, Validation loss: 0.32152456045150757


Epoch 28: Train err: 0.0, Train loss: 0.32152456045150757 | Validation err: 0.043478260869565216, Validation loss: 0.29982027411460876

Epoch 29: Train err: 0.043478260869565216, Train loss: 0.29982027411460876 | Validation err: 0.0, Validation loss: 0.2293560653924942

Epoch 30: Train err: 0.0, Train loss: 0.2293560653924942 | Validation err: 0.021739130434782608, Validation loss: 0.18350203335285187

Finished Training

```
In [10]: model_path = get_model_name("Overfit", batch_size=46, learning_rate=0.001, epoch=29)
         plot_training_curve(model_path)
```



### 3. Hyperparameter Search [10 pt]

#### Part (a) - 1 pt

List 3 hyperparameters that you think are most worth tuning. Choose at least one hyperparameter related to the model architecture.

- Learning Rate
- Batch Size
- Amount of layers (We need to go deeper!)

#### Part (b) - 5 pt

Tune the hyperparameters you listed in Part (a), trying as many values as you need to until you feel satisfied that you are getting a good model. Plot the training curve of at least 4 different hyperparameter settings.

```
In [12]: # Using high Learning rate
model = Gesture("IamSpeed")
if torch.cuda.is_available():
    model = model.cuda()
train(model, learning_rate=0.01)
```

Epoch 1: Train err: 0.8861971830985915, Train loss: 2.653315993150075 | Validation err: 0.900452  
4886877828, Validation loss: 2.1972561180591583

Epoch 2: Train err: 0.895774647887324, Train loss: 2.198790693283081 | Validation err: 0.9004524  
886877828, Validation loss: 2.1975536346435547

Epoch 3: Train err: 0.8918309859154929, Train loss: 2.198551913102468 | Validation err: 0.886877  
8280542986, Validation loss: 2.197644978761673

Epoch 4: Train err: 0.8940845070422535, Train loss: 2.1984989802042643 | Validation err: 0.88687  
78280542986, Validation loss: 2.1976890861988068

Epoch 5: Train err: 0.8935211267605634, Train loss: 2.198500128587087 | Validation err: 0.886877  
8280542986, Validation loss: 2.197710156440735

Epoch 6: Train err: 0.8935211267605634, Train loss: 2.198524598280589 | Validation err: 0.886877  
8280542986, Validation loss: 2.19770884513855

Epoch 7: Train err: 0.8935211267605634, Train loss: 2.1985485235850017 | Validation err: 0.88687  
78280542986, Validation loss: 2.197702467441559

Epoch 8: Train err: 0.8935211267605634, Train loss: 2.1985684593518573 | Validation err: 0.88687  
78280542986, Validation loss: 2.197695314884186

Epoch 9: Train err: 0.8935211267605634, Train loss: 2.1985846757888794 | Validation err: 0.88687  
78280542986, Validation loss: 2.197688966989517

Epoch 10: Train err: 0.8935211267605634, Train loss: 2.1985979715983075 | Validation err: 0.8868  
778280542986, Validation loss: 2.197684019804001

Epoch 11: Train err: 0.8935211267605634, Train loss: 2.1986090421676634 | Validation err: 0.8868  
778280542986, Validation loss: 2.1976794600486755

Epoch 12: Train err: 0.8935211267605634, Train loss: 2.1986183961232504 | Validation err: 0.8868  
778280542986, Validation loss: 2.1976757049560547

Epoch 13: Train err: 0.8935211267605634, Train loss: 2.1986263434092206 | Validation err: 0.8868  
778280542986, Validation loss: 2.197672575712204

Epoch 14: Train err: 0.8935211267605634, Train loss: 2.1986332337061563 | Validation err: 0.8868  
778280542986, Validation loss: 2.1976699233055115

Epoch 15: Train err: 0.8935211267605634, Train loss: 2.198639174302419 | Validation err: 0.88687  
78280542986, Validation loss: 2.197667568922043

Epoch 16: Train err: 0.8935211267605634, Train loss: 2.1986443758010865 | Validation err: 0.8868  
778280542986, Validation loss: 2.1976655423641205

Epoch 17: Train err: 0.8935211267605634, Train loss: 2.198648937543233 | Validation err: 0.88687  
78280542986, Validation loss: 2.1976635456085205

Epoch 18: Train err: 0.8935211267605634, Train loss: 2.1986531019210815 | Validation err: 0.8868  
778280542986, Validation loss: 2.1976619362831116

Epoch 19: Train err: 0.8935211267605634, Train loss: 2.198656709988912 | Validation err: 0.88687  
78280542986, Validation loss: 2.1976605653762817

Epoch 20: Train err: 0.8935211267605634, Train loss: 2.1986599365870156 | Validation err: 0.8868  
778280542986, Validation loss: 2.1976593136787415

Epoch 21: Train err: 0.8935211267605634, Train loss: 2.19866289695104 | Validation err: 0.886877  
8280542986, Validation loss: 2.1976580917835236

Epoch 22: Train err: 0.8935211267605634, Train loss: 2.198665503660838 | Validation err: 0.88687  
78280542986, Validation loss: 2.1976568400859833

Epoch 23: Train err: 0.8935211267605634, Train loss: 2.198667903741201 | Validation err: 0.88687  
78280542986, Validation loss: 2.1976559460163116

Epoch 24: Train err: 0.8935211267605634, Train loss: 2.1986700534820556 | Validation err: 0.8868  
778280542986, Validation loss: 2.197655111551285

Epoch 25: Train err: 0.8935211267605634, Train loss: 2.198672036329905 | Validation err: 0.88687  
78280542986, Validation loss: 2.1976543366909027

Epoch 26: Train err: 0.8935211267605634, Train loss: 2.19867384036382 | Validation err: 0.886877  
8280542986, Validation loss: 2.19765368103981

Epoch 27: Train err: 0.8935211267605634, Train loss: 2.198675497372945 | Validation err: 0.88687  
78280542986, Validation loss: 2.19765305519104



Epoch 28: Train err: 0.8935211267605634, Train loss: 2.198677043120066 | Validation err: 0.88687  
78280542986, Validation loss: 2.1976523995399475

Epoch 29: Train err: 0.8935211267605634, Train loss: 2.1986783822377522 | Validation err: 0.8868  
778280542986, Validation loss: 2.197651892900467

Epoch 30: Train err: 0.8935211267605634, Train loss: 2.1986796657244363 | Validation err: 0.8868  
778280542986, Validation loss: 2.197651356458664

Finished Training

```
In [13]: model_path = get_model_name("IamSpeed", batch_size=30, learning_rate=0.01, epoch=29)
          plot_training_curve(model_path)
```



```
In [14]: # Using small Learning rate
model = Gesture("IamSlow")
if torch.cuda.is_available():
    model = model.cuda()
train(model, learning_rate=0.0001)
```

Epoch 1: Train err: 0.7397183098591549, Train loss: 2.089482625325521 | Validation err: 0.6108597285067874, Validation loss: 1.80134916305542

Epoch 2: Train err: 0.5397183098591549, Train loss: 1.5850840985774994 | Validation err: 0.4072398190045249, Validation loss: 1.2967552542686462

Epoch 3: Train err: 0.39774647887323944, Train loss: 1.2496534526348113 | Validation err: 0.33031674208144796, Validation loss: 1.0944920480251312

Epoch 4: Train err: 0.3340845070422535, Train loss: 1.0848199675480525 | Validation err: 0.31221719457013575, Validation loss: 1.0153076574206352

Epoch 5: Train err: 0.3104225352112676, Train loss: 0.9822133779525757 | Validation err: 0.31221719457013575, Validation loss: 0.9776335209608078

Epoch 6: Train err: 0.29070422535211266, Train loss: 0.9092458556095759 | Validation err: 0.29411764705882354, Validation loss: 0.9581030979752541

Epoch 7: Train err: 0.27098591549295775, Train loss: 0.850671782096227 | Validation err: 0.28054298642533937, Validation loss: 0.9451870694756508

Epoch 8: Train err: 0.2563380281690141, Train loss: 0.801503864924113 | Validation err: 0.27601809954751133, Validation loss: 0.9300251342356205

Epoch 9: Train err: 0.24845070422535212, Train loss: 0.7565458506345749 | Validation err: 0.27149321266968324, Validation loss: 0.9195008017122746

Epoch 10: Train err: 0.23380281690140844, Train loss: 0.7151138871908188 | Validation err: 0.27149321266968324, Validation loss: 0.9121885187923908

Epoch 11: Train err: 0.22422535211267605, Train loss: 0.6754198268055915 | Validation err: 0.26244343891402716, Validation loss: 0.8999729380011559

Epoch 12: Train err: 0.2123943661971831, Train loss: 0.6372159322102865 | Validation err: 0.2579185520361991, Validation loss: 0.8927652910351753

Epoch 13: Train err: 0.19830985915492957, Train loss: 0.6017636880278587 | Validation err: 0.248868778280543, Validation loss: 0.8902707509696484

Epoch 14: Train err: 0.1808450704225352, Train loss: 0.5671152859926224 | Validation err: 0.248868778280543, Validation loss: 0.8895285502076149

Epoch 15: Train err: 0.1667605633802817, Train loss: 0.5341965203483899 | Validation err: 0.24434389140271492, Validation loss: 0.8874035775661469

Epoch 16: Train err: 0.1543661971830986, Train loss: 0.5024464497963588 | Validation err: 0.23529411764705882, Validation loss: 0.8833070397377014

Epoch 17: Train err: 0.14535211267605633, Train loss: 0.4717484273016453 | Validation err: 0.22624434389140272, Validation loss: 0.8874070160090923

Epoch 18: Train err: 0.13859154929577464, Train loss: 0.4415501577158769 | Validation err: 0.23529411764705882, Validation loss: 0.8924807906150818

Epoch 19: Train err: 0.1295774647887324, Train loss: 0.41306864445408187 | Validation err: 0.248868778280543, Validation loss: 0.8946541845798492

Epoch 20: Train err: 0.11830985915492957, Train loss: 0.38640020253757634 | Validation err: 0.24434389140271492, Validation loss: 0.9005938209593296

Epoch 21: Train err: 0.10197183098591549, Train loss: 0.36006149016320704 | Validation err: 0.23076923076923078, Validation loss: 0.904574379324913

Epoch 22: Train err: 0.09352112676056339, Train loss: 0.334469689677159 | Validation err: 0.22624434389140272, Validation loss: 0.9092593491077423

Epoch 23: Train err: 0.08394366197183098, Train loss: 0.30943756687144436 | Validation err: 0.23529411764705882, Validation loss: 0.9118237048387527

Epoch 24: Train err: 0.07661971830985916, Train loss: 0.28565388011435666 | Validation err: 0.23076923076923078, Validation loss: 0.9194266386330128

Epoch 25: Train err: 0.06535211267605634, Train loss: 0.2631293329099814 | Validation err: 0.24434389140271492, Validation loss: 0.9325841702520847

Epoch 26: Train err: 0.06253521126760564, Train loss: 0.24215672469387453 | Validation err: 0.23529411764705882, Validation loss: 0.947666946798563

Epoch 27: Train err: 0.05802816901408451, Train loss: 0.22274858175466458 | Validation err: 0.24434389140271492, Validation loss: 0.9617388695478439


Epoch 28: Train err: 0.0523943661971831, Train loss: 0.20510432831943035 | Validation err: 0.248868778280543, Validation loss: 0.9857555069029331

Epoch 29: Train err: 0.044507042253521124, Train loss: 0.18863991120209297 | Validation err: 0.2579185520361991, Validation loss: 0.99337213113904

Epoch 30: Train err: 0.038873239436619716, Train loss: 0.17585000383357208 | Validation err: 0.26244343891402716, Validation loss: 1.0103040263056755

Finished Training

```
In [15]: model_path = get_model_name("IamSlow", batch_size=30, learning_rate=0.0001, epoch=29)
         plot_training_curve(model_path)
```



```
In [17]: # Using smaller batch size
model = Gesture("IamNiche")
if torch.cuda.is_available():
    model = model.cuda()
train(model, learning_rate=0.001, batch_size=10)
```



Epoch 1: Train err: 0.5588732394366197, Train loss: 1.6038683247700165 | Validation err: 0.3574660633484163, Validation loss: 0.9542100662770479

Epoch 2: Train err: 0.3335211267605634, Train loss: 0.9839736539661215 | Validation err: 0.3031674208144796, Validation loss: 0.8721595665682917

Epoch 3: Train err: 0.2704225352112676, Train loss: 0.8247075809186764 | Validation err: 0.29411764705882354, Validation loss: 0.8753649592399597

Epoch 4: Train err: 0.20788732394366197, Train loss: 0.6420469575802262 | Validation err: 0.27601809954751133, Validation loss: 0.8883483967055446

Epoch 5: Train err: 0.16901408450704225, Train loss: 0.506746675693587 | Validation err: 0.2579185520361991, Validation loss: 0.8410679847002029

Epoch 6: Train err: 0.12845070422535212, Train loss: 0.3892295352612319 | Validation err: 0.23076923076923078, Validation loss: 0.8243717107435932

Epoch 7: Train err: 0.10535211267605633, Train loss: 0.31452892173808905 | Validation err: 0.248868778280543, Validation loss: 0.9585151086034982

Epoch 8: Train err: 0.08732394366197183, Train loss: 0.25921955714891715 | Validation err: 0.23529411764705882, Validation loss: 0.9190404159826753

Epoch 9: Train err: 0.06816901408450704, Train loss: 0.21483508553460576 | Validation err: 0.22171945701357465, Validation loss: 0.9573011669455572

Epoch 10: Train err: 0.05971830985915493, Train loss: 0.1904104334351952 | Validation err: 0.2669683257918552, Validation loss: 0.940740624928604

Epoch 11: Train err: 0.036619718309859155, Train loss: 0.11312431484602099 | Validation err: 0.20361990950226244, Validation loss: 0.867730961289572

Epoch 12: Train err: 0.023098591549295774, Train loss: 0.07192807377079304 | Validation err: 0.22171945701357465, Validation loss: 0.9445845984629831

Epoch 13: Train err: 0.03830985915492958, Train loss: 0.11929236457331938 | Validation err: 0.27601809954751133, Validation loss: 1.3667503940341506

Epoch 14: Train err: 0.024788732394366197, Train loss: 0.09508363564171386 | Validation err: 0.22624434389140272, Validation loss: 1.0841436822412631

Epoch 15: Train err: 0.015211267605633802, Train loss: 0.04133176174766274 | Validation err: 0.2171945701357466, Validation loss: 1.1720354431511268

Epoch 16: Train err: 0.009577464788732394, Train loss: 0.03949220019243962 | Validation err: 0.23076923076923078, Validation loss: 1.5628735855545686

Epoch 17: Train err: 0.0028169014084507044, Train loss: 0.009453807230650025 | Validation err: 0.19457013574660634, Validation loss: 1.4289549893355402

Epoch 18: Train err: 0.0, Train loss: 0.0031541401207309822 | Validation err: 0.17647058823529413, Validation loss: 1.2303220608483552

Epoch 19: Train err: 0.0, Train loss: 0.0010509374581461547 | Validation err: 0.19004524886877827, Validation loss: 1.2487919085211694

Epoch 20: Train err: 0.0, Train loss: 0.0005100881446175505 | Validation err: 0.19004524886877827, Validation loss: 1.261064467060746

Epoch 21: Train err: 0.0, Train loss: 0.0003401879156380285 | Validation err: 0.19004524886877827, Validation loss: 1.2864859074672752

Epoch 22: Train err: 0.0, Train loss: 0.00025972539161225227 | Validation err: 0.18099547511312217, Validation loss: 1.3069774570606412

Epoch 23: Train err: 0.0, Train loss: 0.00021257766785525705 | Validation err: 0.18099547511312217, Validation loss: 1.3232316346155615

Epoch 24: Train err: 0.0, Train loss: 0.00018001613685016983 | Validation err: 0.18099547511312217, Validation loss: 1.3389510469214954

Epoch 25: Train err: 0.0, Train loss: 0.0001547368175876473 | Validation err: 0.18552036199095023, Validation loss: 1.3525881379647815

Epoch 26: Train err: 0.0, Train loss: 0.0001349771955629484 | Validation err: 0.18099547511312217, Validation loss: 1.366693511075062

Epoch 27: Train err: 0.0, Train loss: 0.00011886167661238728 | Validation err: 0.18099547511312217, Validation loss: 1.380254660206642

Epoch 28: Train err: 0.0, Train loss: 0.00010513397822521501 | Validation err: 0.18099547511312217, Validation loss: 1.3931754139105197

Epoch 29: Train err: 0.0, Train loss: 9.33245063400075e-05 | Validation err: 0.18099547511312217, Validation loss: 1.4068815842594777

Epoch 30: Train err: 0.0, Train loss: 8.324854170546407e-05 | Validation err: 0.18099547511312217, Validation loss: 1.41935036576773

Finished Training

```
In [18]: model_path = get_model_name("IamNiche", batch_size=10, learning_rate=0.001, epoch=29)
         plot_training_curve(model_path)
```



In [19]: *# Try out a different architecture with deeper layers*

```
class DeepGesture(nn.Module):
    def __init__(self, name):
        super(DeepGesture, self).__init__()
        self.name = name
        self.conv1 = nn.Conv2d(3, 10, 5)
        self.conv2 = nn.Conv2d(10, 20, 3)
        self.conv3 = nn.Conv2d(20, 40, 3)

        self.pool = nn.MaxPool2d(2, 2)

        self.fc1 = nn.Linear(40 * 26 * 26, 224)
        self.fc2 = nn.Linear(224, 9) # ABCDEFGHI

    def forward(self, x):
        x = self.pool(F.relu(self.conv1(x)))
        x = self.pool(F.relu(self.conv2(x)))
        x = self.pool(F.relu(self.conv3(x)))
        x = x.view(-1, 40 * 26 * 26)
        x = F.relu(self.fc1(x))
        x = self.fc2(x)
        return x

model = DeepGesture("IamDeep")
if torch.cuda.is_available():
    model = model.cuda()
train(model, learning_rate=0.0001)
```



Epoch 1: Train err: 0.856338028169014, Train loss: 2.1845354795455934 | Validation err: 0.7058823529411765, Validation loss: 2.1191424429416656

Epoch 2: Train err: 0.643943661971831, Train loss: 1.8785146017869314 | Validation err: 0.4796380090497738, Validation loss: 1.4971020072698593

Epoch 3: Train err: 0.45577464788732397, Train loss: 1.3559428920348486 | Validation err: 0.36199095022624433, Validation loss: 1.140036553144455

Epoch 4: Train err: 0.3526760563380282, Train loss: 1.1086917221546173 | Validation err: 0.31221719457013575, Validation loss: 1.0222617462277412

Epoch 5: Train err: 0.3211267605633803, Train loss: 0.985493556658427 | Validation err: 0.2986425339366516, Validation loss: 0.9804055616259575

Epoch 6: Train err: 0.29971830985915493, Train loss: 0.908753048380216 | Validation err: 0.2986425339366516, Validation loss: 0.9664740189909935

Epoch 7: Train err: 0.27774647887323944, Train loss: 0.847885578374068 | Validation err: 0.2850678733031674, Validation loss: 0.9463155120611191

Epoch 8: Train err: 0.26028169014084507, Train loss: 0.7990952571233113 | Validation err: 0.2850678733031674, Validation loss: 0.9349631294608116

Epoch 9: Train err: 0.24901408450704227, Train loss: 0.7566410670677821 | Validation err: 0.2850678733031674, Validation loss: 0.9261490181088448

Epoch 10: Train err: 0.23098591549295774, Train loss: 0.7177986587087314 | Validation err: 0.2850678733031674, Validation loss: 0.9218388423323631

Epoch 11: Train err: 0.21746478873239436, Train loss: 0.6806270410617192 | Validation err: 0.25339366515837103, Validation loss: 0.9198045544326305

Epoch 12: Train err: 0.2112676056338028, Train loss: 0.6449140648047129 | Validation err: 0.248868778280543, Validation loss: 0.9226311519742012

Epoch 13: Train err: 0.19549295774647887, Train loss: 0.6102942059437434 | Validation err: 0.24434389140271492, Validation loss: 0.9218592904508114

Epoch 14: Train err: 0.18591549295774648, Train loss: 0.5770264692604542 | Validation err: 0.2398190045248869, Validation loss: 0.9197849631309509

Epoch 15: Train err: 0.1712676056338028, Train loss: 0.5449417454500993 | Validation err: 0.248868778280543, Validation loss: 0.9266286641359329

Epoch 16: Train err: 0.1616901408450704, Train loss: 0.5128569071491559 | Validation err: 0.23529411764705882, Validation loss: 0.9217727184295654

Epoch 17: Train err: 0.14985915492957746, Train loss: 0.4810715968410174 | Validation err: 0.2398190045248869, Validation loss: 0.9244362488389015

Epoch 18: Train err: 0.13408450704225353, Train loss: 0.44897813697655997 | Validation err: 0.23529411764705882, Validation loss: 0.9163704141974449

Epoch 19: Train err: 0.1267605633802817, Train loss: 0.4199917189776897 | Validation err: 0.23076923076923078, Validation loss: 0.9124659150838852

Epoch 20: Train err: 0.11549295774647887, Train loss: 0.3897153403609991 | Validation err: 0.22624434389140272, Validation loss: 0.9107358083128929

Epoch 21: Train err: 0.10873239436619718, Train loss: 0.36040155626833437 | Validation err: 0.22171945701357465, Validation loss: 0.9170565493404865

Epoch 22: Train err: 0.09577464788732394, Train loss: 0.33055752428869406 | Validation err: 0.2081447963800905, Validation loss: 0.935352448374033

Epoch 23: Train err: 0.0856338028169014, Train loss: 0.30267224833369255 | Validation err: 0.21266968325791855, Validation loss: 0.9595853686332703

Epoch 24: Train err: 0.07267605633802816, Train loss: 0.2749189084395766 | Validation err: 0.2171945701357466, Validation loss: 0.9787172898650169

Epoch 25: Train err: 0.06535211267605634, Train loss: 0.24870025012642144 | Validation err: 0.2171945701357466, Validation loss: 1.0048662573099136

Epoch 26: Train err: 0.05295774647887324, Train loss: 0.22200418865929047 | Validation err: 0.2081447963800905, Validation loss: 1.034901313483715

Epoch 27: Train err: 0.04732394366197183, Train loss: 0.19922662259389956 | Validation err: 0.21266968325791855, Validation loss: 1.0672782063484192



Epoch 28: Train err: 0.04056338028169014, Train loss: 0.177862517721951 | Validation err: 0.2171945701357466, Validation loss: 1.1108272522687912

Epoch 29: Train err: 0.03436619718309859, Train loss: 0.15950245624408127 | Validation err: 0.22171945701357465, Validation loss: 1.1500873267650604

Epoch 30: Train err: 0.03436619718309859, Train loss: 0.14405676880851387 | Validation err: 0.2171945701357466, Validation loss: 1.213636301457882

Finished Training

```
In [20]: model_path = get_model_name("IamDeep", batch_size=30, learning_rate=0.0001, epoch=29)
         plot_training_curve(model_path)
```



## Part (c) - 2 pt

Choose the best model out of all the ones that you have trained. Justify your choice.

All models we tried seem to have high Validation Loss over time, which is a sign of Overfitting (bad). However, between all of them, I would choose the last model (Deeper network, batch\_size=30, learning\_rate=0.0001, num\_epoch=30) because its Validation error reaches the lowest value AND is flattened (consistent) overtime.

## Part (d) - 2 pt

Report the test accuracy of your best model. You should only do this step once and prior to this step you should have only used the training and validation data.

```
In [21]: criterion = nn.CrossEntropyLoss()
test_err, test_loss = evaluate(model, test_loader, criterion)

print("Test error:", test_err)
print("Test loss:", test_loss)
```

Test error: 0.10762331838565023  
Test loss: 0.42889735606488305

## 4. Transfer Learning [15 pt]

For many image classification tasks, it is generally not a good idea to train a very large deep neural network model from scratch due to the enormous compute requirements and lack of sufficient amounts of training data.

One of the better options is to try using an existing model that performs a similar task to the one you need to solve. This method of utilizing a pre-trained network for other similar tasks is broadly termed **Transfer Learning**. In this assignment, we will use Transfer Learning to extract features from the hand gesture images. Then, train a smaller network to use these features as input and classify the hand gestures.

As you have learned from the CNN lecture, convolution layers extract various features from the images which get utilized by the fully connected layers for correct classification. AlexNet architecture played a pivotal role in establishing Deep Neural Nets as a go-to tool for image classification problems and we will use an ImageNet pre-trained AlexNet model to extract features in this assignment.

## Part (a) - 5 pt

Here is the code to load the AlexNet network, with pretrained weights. When you first run the code, PyTorch will download the pretrained weights from the internet.

```
In [28]: import torchvision.models
alexnet = torchvision.models.alexnet(pretrained=True)
```

The alexnet model is split up into two components: *alexnet.features* and *alexnet.classifier*. The first neural network component, *alexnet.features*, is used to compute convolutional features, which are taken as input in *alexnet.classifier*.

The neural network alexnet.features expects an image tensor of shape Nx3x224x224 as input and it will output a tensor of shape Nx256x6x6 . (N = batch size).

Compute the AlexNet features for each of your training, validation, and test data. Here is an example code snippet showing how you can compute the AlexNet features for some images (your actual code might be different):

**Save the computed features.** You will be using these features as input to your neural network in Part (b), and you do not want to re-compute the features every time. Instead, run *alexnet.features* once for each image, and save the result.

```
In [43]: import os

# guaranteed that overfit set is separated from training and validation
transfer_train_loader, transfer_val_loader, transfer_test_loader, transfer_overfit_loader = get_

training_store = []
validation_store = []
testing_store = []
overfit_store = []

for img, label in transfer_train_loader:
    features = alexnet.features(img)
    features_tensor = torch.from_numpy(features.detach().numpy()).squeeze(0)
    training_store.append((features_tensor, label))

for img, label in transfer_val_loader:
```

```

features = alexnet.features(img)
features_tensor = torch.from_numpy(features.detach().numpy()).squeeze(0)
validation_store.append((features_tensor, label))

for img, label in transfer_test_loader:
    features = alexnet.features(img)
    features_tensor = torch.from_numpy(features.detach().numpy()).squeeze(0)
    testing_store.append((features_tensor, label))

for img, label in transfer_overfit_loader:
    features = alexnet.features(img)
    features_tensor = torch.from_numpy(features.detach().numpy()).squeeze(0)
    overfit_store.append((features_tensor, label))

```

## Part (b) - 3 pt

Build a convolutional neural network model that takes as input these AlexNet features, and makes a prediction. Your model should be a subclass of nn.Module.

Explain your choice of neural network architecture: how many layers did you choose? What types of layers did you use: fully-connected or convolutional? What about other decisions like pooling layers, activation functions, number of channels / hidden units in each layer?

Here is an example of how your model may be called:

```

In [44]: # Only need classification layer since all CNN feature extractions are done through AlexNet
# Chose FC layer size based on original AlexNet design
# Activation: ReLU

class AlexNet(nn.Module):
    def __init__(self, name):
        super(AlexNet, self).__init__()
        self.name = name
        self.fc1 = nn.Linear(256 * 6 * 6, 32)
        self.fc2 = nn.Linear(32, 9)

    def forward(self, x):
        x = x.view(-1, 256 * 6 * 6)
        x = F.relu(self.fc1(x))
        x = self.fc2(x)
        return x

```

## Part (c) - 5 pt

Train your new network, including any hyperparameter tuning. Plot and submit the training curve of your best model only.

Note: Depending on how you are caching (saving) your AlexNet features, PyTorch might still be tracking updates to the **AlexNet weights**, which we are not tuning. One workaround is to convert your AlexNet feature tensor into a numpy array, and then back into a PyTorch tensor.

```

In [45]: model = AlexNet("Transfer")
if torch.cuda.is_available():
    model = model.cuda()
train(model, overfitting=True, batch_size=30, learning_rate=0.0001, num_epochs=30, replace=[trai

```

WARNING: only overfit testing should be here

WARNING: only transfer learning should be here

Epoch 1: Train err: 0.3036619718309859, Train loss: 1.0847654354876615 | Validation err: 0.2613065326633166, Validation loss: 0.7657881107312351

Epoch 2: Train err: 0.09521126760563381, Train loss: 0.42853839818142275 | Validation err: 0.1507537688442211, Validation loss: 0.5529330226232358

Epoch 3: Train err: 0.056901408450704224, Train loss: 0.2489131877365732 | Validation err: 0.09045226130653267, Validation loss: 0.45229160907179883

Epoch 4: Train err: 0.029295774647887324, Train loss: 0.15597657999262313 | Validation err: 0.08542713567839195, Validation loss: 0.39956174211241446

Epoch 5: Train err: 0.013521126760563381, Train loss: 0.10176572216012779 | Validation err: 0.08542713567839195, Validation loss: 0.36615841644703123

Epoch 6: Train err: 0.009014084507042254, Train loss: 0.06742422185551017 | Validation err: 0.08542713567839195, Validation loss: 0.33652714205205125

Epoch 7: Train err: 0.0033802816901408453, Train loss: 0.04473412071367596 | Validation err: 0.08040201005025126, Validation loss: 0.31684677446842663

Epoch 8: Train err: 0.0016901408450704226, Train loss: 0.02920821819117132 | Validation err: 0.07537688442211055, Validation loss: 0.3036408945719775

Epoch 9: Train err: 0.0011267605633802818, Train loss: 0.01914534719272106 | Validation err: 0.07537688442211055, Validation loss: 0.3018788439733779

Epoch 10: Train err: 0.0005633802816901409, Train loss: 0.012600966079339125 | Validation err: 0.07537688442211055, Validation loss: 0.3052571710255885

Epoch 11: Train err: 0.0005633802816901409, Train loss: 0.009003958460711952 | Validation err: 0.07537688442211055, Validation loss: 0.31066260782775995

Epoch 12: Train err: 0.0011267605633802818, Train loss: 0.007256976771134037 | Validation err: 0.06532663316582915, Validation loss: 0.32115360193244546

Epoch 13: Train err: 0.0011267605633802818, Train loss: 0.006124770340234892 | Validation err: 0.06532663316582915, Validation loss: 0.3342275365027025

Epoch 14: Train err: 0.0011267605633802818, Train loss: 0.005240989505440927 | Validation err: 0.06532663316582915, Validation loss: 0.3513627858500299

Epoch 15: Train err: 0.0011267605633802818, Train loss: 0.005598771914923018 | Validation err: 0.06532663316582915, Validation loss: 0.36940493418324905

Epoch 16: Train err: 0.0011267605633802818, Train loss: 0.005582923383034132 | Validation err: 0.06532663316582915, Validation loss: 0.39580210003739835

Epoch 17: Train err: 0.0011267605633802818, Train loss: 0.004456707582628901 | Validation err: 0.06532663316582915, Validation loss: 0.403875485935086

Epoch 18: Train err: 0.0011267605633802818, Train loss: 0.006822154298080502 | Validation err: 0.06532663316582915, Validation loss: 0.4341712613727664

Epoch 19: Train err: 0.0011267605633802818, Train loss: 0.006179120398283187 | Validation err: 0.06532663316582915, Validation loss: 0.45454105023543706

Epoch 20: Train err: 0.0011267605633802818, Train loss: 0.006264986240026428 | Validation err: 0.07035175879396985, Validation loss: 0.47560062157676053

Epoch 21: Train err: 0.0011267605633802818, Train loss: 0.006532649485183593 | Validation err: 0.07035175879396985, Validation loss: 0.5013516043531382

Epoch 22: Train err: 0.0005633802816901409, Train loss: 0.005773988908090884 | Validation err: 0.07035175879396985, Validation loss: 0.508955570331471

Epoch 23: Train err: 0.0011267605633802818, Train loss: 0.005479529750946176 | Validation err: 0.06532663316582915, Validation loss: 0.5258917389151211

Epoch 24: Train err: 0.0011267605633802818, Train loss: 0.004958175305614107 | Validation err: 0.06532663316582915, Validation loss: 0.544205167477745

Epoch 25: Train err: 0.0011267605633802818, Train loss: 0.004968356897384879 | Validation err: 0.06532663316582915, Validation loss: 0.5736909306675745

Epoch 26: Train err: 0.0011267605633802818, Train loss: 0.007387196722680551 | Validation err: 0.07537688442211055, Validation loss: 0.633221154507613

Epoch 27: Train err: 0.0011267605633802818, Train loss: 0.00708199438816879 | Validation err: 0.07537688442211055, Validation loss: 0.6539957123562324



Epoch 28: Train err: 0.0011267605633802818, Train loss: 0.007091564141300404 | Validation err: 0.08040201005025126, Validation loss: 0.6987219198467014

Epoch 29: Train err: 0.0005633802816901409, Train loss: 0.006157495049453532 | Validation err: 0.07035175879396985, Validation loss: 0.653696787518075

Epoch 30: Train err: 0.0011267605633802818, Train loss: 0.005313563793034858 | Validation err: 0.07035175879396985, Validation loss: 0.675511242401347

Finished Training

```
In [46]: model_path = get_model_name("Transfer", batch_size=30, learning_rate=0.0001, epoch=29)
         plot_training_curve(model_path)
```



## Part (d) - 2 pt

Report the test accuracy of your best model. How does the test accuracy compare to Part 3(d) without transfer learning?

```
In [47]: criterion = nn.CrossEntropyLoss()
         test_err, test_loss = evaluate(model, testing_store, criterion)

         print("Test error:", test_err)
         print("Test loss:", test_loss)
```

Test error: 0.04522613065326633  
Test loss: 0.20237362812151202

As we can see, this model (**4.5% error, 20.2% loss**) performs better compared to Part3(d) (**10% error, 42.8% loss**)

## 5. Additional Testing [5 pt]

For this question you should just reuse the small dataset that you used to answer question 2c.

Using the best transfer learning model developed in Part 4. Report the test accuracy on your sample images and how it compares to the test accuracy obtained in Part 4(d)? How well did your model do for the different hand gestures? Provide an explanation for why you think your model performed the way it did?

```
In [48]: criterion = nn.CrossEntropyLoss()
test_err, test_loss = evaluate(model, overfit_store, criterion)

print("Test error:", test_err)
print("Test loss:", test_loss)
```

Test error: 0.06521739130434782

Test loss: 0.17022235944834394

The test accuracy is approximately the same as Part 4(d), only a little bit lower.

- Good accuracy since the AlexNet CNN has deeper layers and pre-trained weights that can extract image features very well.
- Quite similar to testing accuracy since the datasets are guaranteed to be separate (split from `get_data_loader`) so for the model, the test set and overfit set are similar (things that it has never seen before)
- Accuracy is a bit less than testing, also possibly because the number of data tested in this set is much smaller, so errors are more expensive